GPU engines in the 2026-10-25 A100 study versus this H100 campaign (non-equivalent hosts)

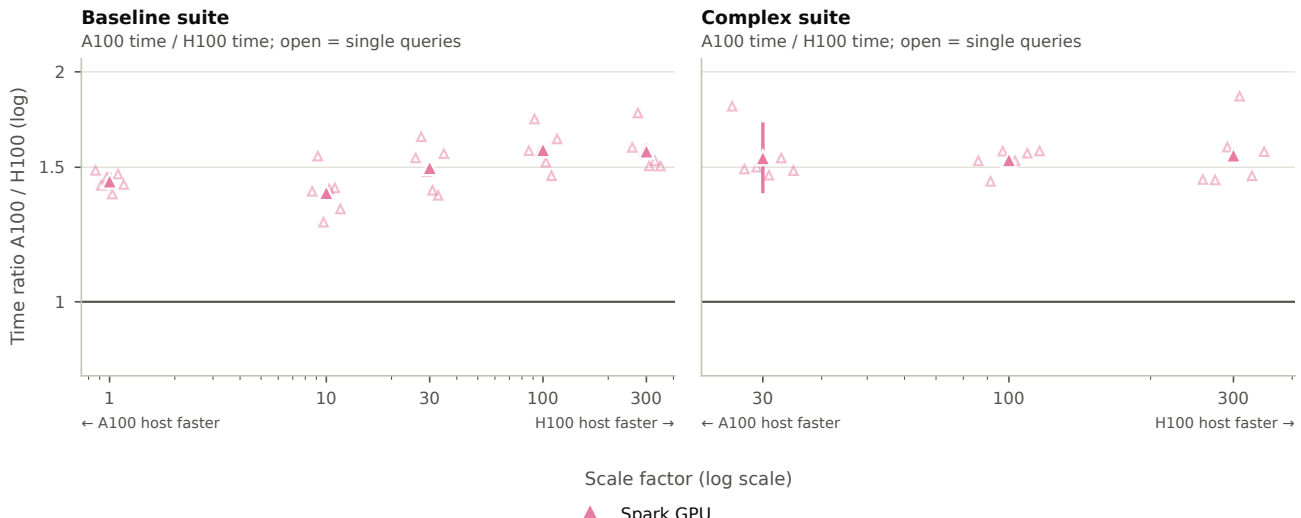

NOT an equivalent comparison: the 2026-10-25 A100 study (saved evidence of a previous study, not the A100 campaign; 1 x NVIDIA A100-SXM4-40GB, 30 vCPU, 216 GiB) versus NVIDIA H100 80GB HBM3, 26 vCPU, 221.2 GiB. Differences: Polars 2.0.0 against 1.44.2 (CPU engine; not in the ratio); 30 against 26 threads; protocol: 1 session per cell against 4 rounds; fixed engine rotation per scale against rotated, balanced orders; query cleanup policy, failure policy, suite version, Spark GPU profile not recorded in the earlier study. Polars GPU was not run in that study. CPU engines are not compared. The earlier study is one session per cell, so the bars show round-to-round variation on the H100 side only.