GPU-generation ladder: suite-level time relative to H100 by engine and host

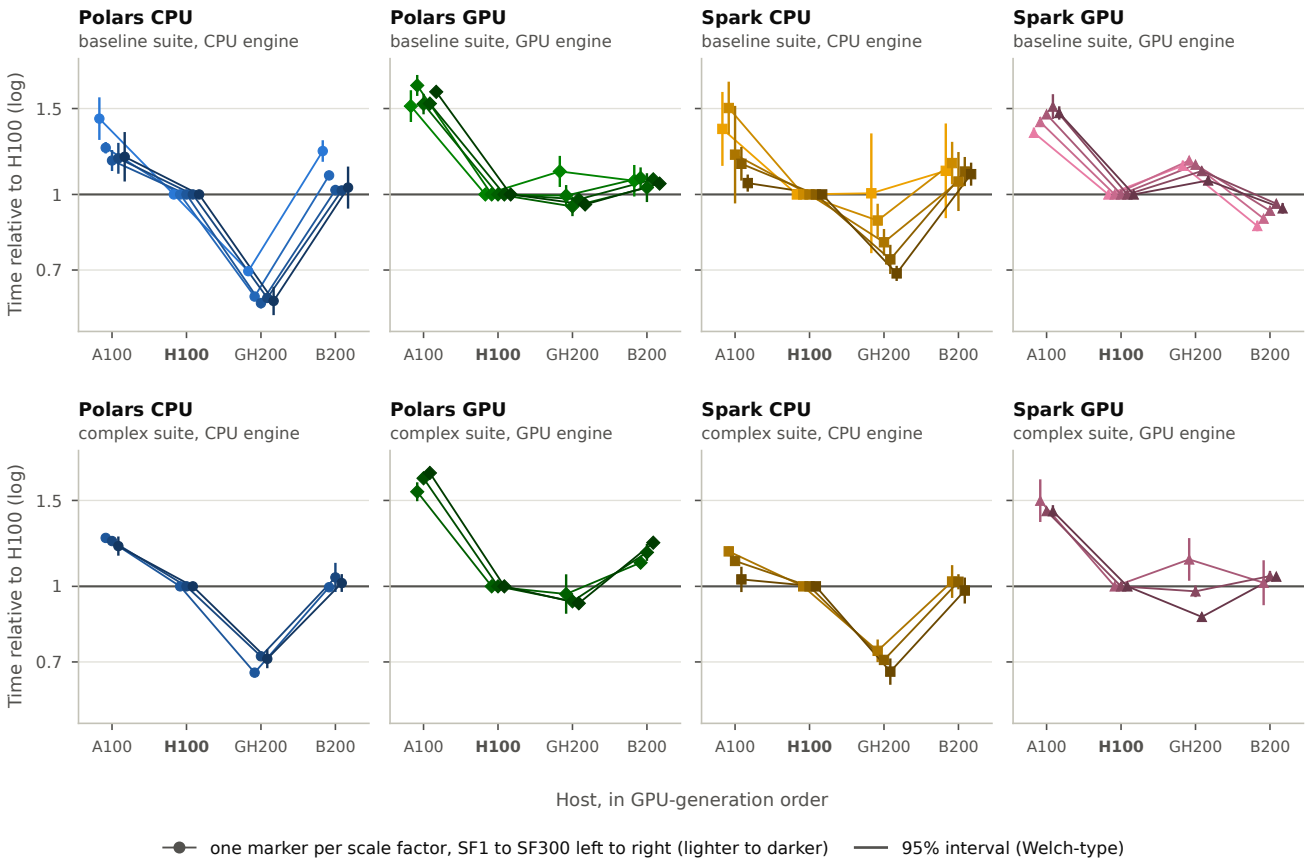

Time on the host divided by time on H100 (geometric mean of round medians over the queries timed in every round on both hosts); below 1 the host is faster. Bars: Welch-type 95% interval, rounds unpaired. Ladder order: A100 (Ampere), H100 (Hopper), GH200 (Grace Hopper), B200 (Blackwell). Each host is a separate campaign; CPU, memory, driver and (GH200) architecture differ, so GPU-engine differences are not GPU-only effects and CPU-engine differences reflect the host CPU.